

1xp9qwl1r

February 20, 2025

```
[3]: # Import all the Dependencies

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
[4]: # Import the heart_csv dataset
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving heart.csv to heart.csv

```
[6]: # Load the dataset into a dataframe
heart_data = pd.read_csv('heart.csv')
```

```
[7]: # Print the first 5 rows of the dataset
heart_data.head()
```

```
[7]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope
ca	thal	target									
0	52	1	0	125	212	0	1	168	0	1.0	2
2	3		0								
1	53	1	0	140	203	1	0	155	1	3.1	0
0	3		0								
2	70	1	0	145	174	0	1	125	1	2.6	0
0	3		0								
3	61	1	0	148	203	0	1	161	0	0.0	2
1	3		0								
4	62	0	0	138	294	1	1	106	0	1.9	1
3	2		0								

```
[8]: # Print the last 5 rows of the dataset
heart_data.tail()
```

```
[8]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope
      ca  thal  target
1020  59    1  1    140    221    0      1    164    1    0.0    2
0     2      1
1021  60    1  0    125    258    0      0    141    1    2.8    1
1     3      0
1022  47    1  0    110    275    0      0    118    1    1.0    1
1     2      0
1023  50    0  0    110    254    0      0    159    0    0.0    2
0     2      1
1024  54    1  0    120    188    0      1    113    0    1.4    1
1     3      0
```

```
[9]: # Number of rows & columns in the dataset
heart_data.shape
```

```
[9]: (1025, 14)
```

```
[10]: # Getting information about the data
heart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
[11]: # Checking for missing values
heart_data.isnull().sum()
```

```
[11]: age      0
      sex      0
      cp       0
      trestbps 0
      chol     0
      fbs      0
      restecg  0
      thalach  0
      exang    0
      oldpeak  0
      slope    0
      ca       0
      thal     0
      target   0
      dtype: int64
```

```
[12]: # Get statistical measure of the data
      heart_data.describe()
```

```
[12]:
```

	age	sex	cp	...	ca	thal
target						
count	1025.000000	1025.000000	1025.000000	...	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	...	0.754146	2.323902
std	9.072290	0.460373	1.029641	...	1.030798	0.620660
min	29.000000	0.000000	0.000000	...	0.000000	0.000000
25%	48.000000	0.000000	0.000000	...	0.000000	2.000000
50%	56.000000	1.000000	1.000000	...	0.000000	2.000000
75%	61.000000	1.000000	2.000000	...	1.000000	3.000000
max	77.000000	1.000000	3.000000	...	4.000000	3.000000

[8 rows x 14 columns]

```
[13]: # How is the target value distributed
      heart_data['target'].value_counts()
```

```
[13]: target
      1    526
      0    499
      Name: count, dtype: int64
```

```
[14]: # Splitting the features and target
X = heart_data.drop(columns='target', axis=1)
Y = heart_data['target']
```

```
[15]: print(X)
print(Y)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope
ca	thal										
0	52	1	0	125	212	0	1	168	0	1.0	2
2	3										
1	53	1	0	140	203	1	0	155	1	3.1	0
0	3										
2	70	1	0	145	174	0	1	125	1	2.6	0
0	3										
3	61	1	0	148	203	0	1	161	0	0.0	2
1	3										
4	62	0	0	138	294	1	1	106	0	1.9	1
3	2										
...
...
1020	59	1	1	140	221	0	1	164	1	0.0	2
0	2										
1021	60	1	0	125	258	0	0	141	1	2.8	1
1	3										
1022	47	1	0	110	275	0	0	118	1	1.0	1
1	2										
1023	50	0	0	110	254	0	0	159	0	0.0	2
0	2										
1024	54	1	0	120	188	0	1	113	0	1.4	1
1	3										

[1025 rows x 13 columns]

```
0      0
1      0
2      0
3      0
4      0
..
1020    1
1021    0
1022    0
1023    1
1024    0
```

Name: target, Length: 1025, dtype: int64

```
[16]: # Split the data into training and testing data
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2,
↳stratify=Y, random_state=2)
```

```
[17]: print(X.shape, X_train.shape, X_test.shape)
```

(1025, 13) (820, 13) (205, 13)

```
[18]: # Model training
# Logistic Regression
model = LogisticRegression()
```

```
[19]: # Training the Logistic Regression model with training data
model.fit(X_train, Y_train)
```

/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
[19]: LogisticRegression()
```

```
[20]: # Model Evaluation
# Accuracy score as evaluation metrics
# Accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
[21]: # Finding accuracy_score of training data
print('Accuracy on training data:',training_data_accuracy)
```

Accuracy on training data: 0.8524390243902439

```
[22]: # Accuracy score on the testing data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
[23]: # Finding the accuracy score of test data
print('Accuracy on test data:', test_data_accuracy)
```

Accuracy on test data: 0.8048780487804879

```
[28]: # Building a predictive system
input_data = (61,1,1,0,149,203,0,1,161,0,0,2,1)

# Change the input data to np.array
input_data_as_numpy_array = np.asarray(input_data)

#Reshape the numpy array as we are predicting for only one instance
input_data_reshape = np.reshape(input_data_as_numpy_array, (1,-1))

prediction = model.predict(input_data_reshape)
print(prediction)

if prediction ==[1]:
    print('The person does not have a heart disease')
else:
    print('The person has heart disease')
```

[0]

The person has heart disease

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LogisticRegression was
fitted with feature names
  warnings.warn(
```